

开源地址 : <https://github.com/Vincent-Huang-2000/DATA301>

感觉不错的话, 可以在 Github 上给个 Star

本文档未完结, 最新修订日期: 2023年6月19日

## 提供的资料

May 30 Lecture 28:50 讲到考试中可以访问的资料

- 所有的 lecture notes
- PySpark 的文档

## 题型

- Knowledge [approx. 40 pts]
  - Multiple Choice, choose best answer
- Analysis [approx. 20 pts]
  - Short answer, write 1-4 sentences
- Skill [approx. 40 pts]
  - Apply an algorithm (i.e. compute answer, possibly by calling functions)
  - Implement an algorithm (i.e. write code)

## 考试主题

### Programming

- Map/Reduce functional programming, Message-Passing, Threads, Locks and Atomics, Work Queues, Schedulers, Streaming, MPI, CUDA, SPARK

### Algorithms

- Divide and Conquer, Map, Reduce, Group By, Union, Intersection, Difference, Matrix-Vector and Matrix-Matrix Multiplication, Hashing, PageRank, Graphs, Leader Election, Consensus

### Systems

- SaaS, PaaS, IaaS, storage and networking architectures, virtual machines and their management, job scheduling, cloud resources, heterogeneous processors

### Data and Scale

- 5 Vs (Variety, Velocity, Volume, Veracity, Value), Decomposition, Distributed Data Structures, Memory Hierarchy, Shared memory, Shared-nothing, distributed file systems, replication, communication cost, complexity theory

# Programming

Map/Reduce functional programming (SPARK)

- lambda
- RDD, parallelize, textFile, map, flatMap, filter, reduce, reduceByKey, groupByKey, sortByKey, join, cogroup, cartesian, collect/take, count, countByKey

Message-Passing, Threads, Locks and Atomics, Work Queues, Schedulers

MPI

- rank, send/recv, broadcast, reduce

CUDA

- Numba (jit / cuda.jit)
- Kernel, block, thread, warp

# Algorithms

Divide and Conquer: data and functional parallelism, pipelining

Map, Reduce, Group By

- Frequent Item Sets, Market Basket Analysis

Union, Intersection, Difference

- distance/similarity measures (Jaccard, Cosine, Euclidean/ $L_2$ Norm)
- Content recommendation

Matrix-Vector and Matrix-Matrix Multiplication

Hashing: Shingling, Min-Hashing

Graphs

- Flow/PageRank: random walk, recursive/iterative, traps, teleports
- Clustering: hierarchical (Girvan-Newman/betweenness, modularity)

Online: AdWords/balance

- Simulation: simple N-Body (game of life)
- Leader Election, Consensus (RAFT)

## Systems

- SaaS, PaaS, IaaS
- storage and networking architectures
- virtual machines and their management
- job scheduling, cloud resources [specifically the Google Cloud Platform you've used], heterogeneous processors
- Hardware/low-level parallelism
  - Functional units
  - Instruction level parallelism
  - Out of order execution
  - Dependencies
- Memory Hierarchy, Shared memory, Distributed memory
- SIMD and GPU vector processing: tradeoffs with CPU – less cache/memory, less control logic, more ALU

## Data and Scale

- 5 Vs (Variety, Velocity, Volume, Veracity, Value)
- Decomposition: block, cyclic
- Replication, distributed file systems
- communication cost, complexity theory
  - Amdahl's Law, Gustafson's Law
  - Weak scalability, Strong scalability
  - Elapsed communication cost
  - "roofline performance" and arithmetic intensity

# 期末样题

## multiple choice question

In Market Basket Analysis we might be given the following data:

B1 = {m, c, b} B2 = {m, p, j}

B3 = {m, b} B4 = {b, c, j}

B5 = {m, p, b} B6 = {m, c, x, y}

B7 = {m, c, b, j} B8 = {b, c}

Which relationship has the least confidence?

Select one:

A.  $p \rightarrow m$

B.  $x \rightarrow y$

C.  $c \rightarrow b$

D.  $m \rightarrow x$

在进行市场篮子分析时，我们需要计算关联规则的置信度。置信度定义为：

$$\text{Confidence}(A \rightarrow B) = \text{support}(A \cup B) / \text{support}(A)$$

在这里，A和B是项集，U代表集合的并集。根据这个定义，我们需要计算每个选项的置信度。这需要先找到每个关系中元素A和B的出现频率，以及它们共同出现的频率。

A.  $p \rightarrow m$

$$\text{Confidence}(p \rightarrow m) = \text{support}(\{p, m\}) / \text{support}(p)$$

我们可以看到，{p, m}在B2和B5中出现，所以 $\text{support}(\{p, m\}) = 2$ 。p在B2、B5中出现，所以 $\text{support}(p) = 2$ 。因此， $\text{Confidence}(p \rightarrow m) = 2 / 2 = 1$ 。

B.  $x \rightarrow y$

$$\text{Confidence}(x \rightarrow y) = \text{support}(\{x, y\}) / \text{support}(x)$$

{x, y}只在B6中出现，所以 $\text{support}(\{x, y\}) = 1$ 。x也只在B6中出现，所以 $\text{support}(x) = 1$ 。因此， $\text{Confidence}(x \rightarrow y) = 1 / 1 = 1$ 。

C.  $c \rightarrow b$

$$\text{Confidence}(c \rightarrow b) = \text{support}(\{c, b\}) / \text{support}(c)$$

{c, b}在B1、B4、B7、B8中出现，所以 $\text{support}(\{c, b\}) = 4$ 。c在B1、B4、B6、B7、B8中出现，所以 $\text{support}(c) = 5$ 。因此， $\text{Confidence}(c \rightarrow b) = 4 / 5 = 0.8$ 。

D.  $m \rightarrow x$

$$\text{Confidence}(m \rightarrow x) = \text{support}(\{m, x\}) / \text{support}(m)$$

{m, x}只在B6中出现，所以 $\text{support}(\{m, x\}) = 1$ 。m在B1、B2、B3、B5、B6、B7中出现，所以 $\text{support}(m) = 6$ 。因此， $\text{Confidence}(m \rightarrow x) = 1 / 6 \approx 0.17$ 。

所以，置信度最低的关系是选项D.  $m \rightarrow x$ 。

## short answer question

Web server log files contain records of which page on a server was requested by a user. For example, this is a sample line from a web server log:

```
piweba3y.prodigy.com - 807301196 GET  
/shuttle/missions/missions.html 200 8677
```

Explain how this is example of Velocity (one of the 5 V's of Big Data)

Explain how this is example of Variety (one of the 5 V's of Big Data)

在大数据的五个V中，Velocity(速度)和Variety(多样性)是非常重要的概念。

1. Velocity(速度): Velocity 在大数据中主要指的是数据的生成速度或处理速度。在你提供的例子中，web服务器日志文件是一个很好的展示速度的例子。每当用户请求服务器上的页面时，服务器都会立即生成一条记录。这可能会导致每天，甚至每小时，都会生成大量的新数据。这种高速率的数据生成需要高效的处理和分析方法，这就是大数据速度的一个典型示例。

2. Variety(多样性): Variety 在大数据中是指各种各样的数据类型和来源。在你的例子中，web服务器日志文件也是一个很好的展示多样性的例子。日志文件可能包含各种类型的信息，如请求的IP地址、时间戳、请求的方法(GET或POST)、请求的资源URL、HTTP响应代码，以及响应的大小等。这些数据可能是结构化的(例如，时间戳或响应代码)，也可能是半结构化的(例如，URL)。这些不同类型的数据需要不同的工具和技术来处理和分析，这就是大数据多样性的一个典型示例。

## apply algorithm question

The Jaccard similarity of sets S and T is  $|S \cap T| / |S \cup T|$

Choose the names of two of your friends. Compute the Jaccard similarity between their names where S and T are the sets of letters in each of their names.

假设我有两个朋友，他们的名字是 "Alice" 和 "Elsa"。我们将他们的名字中的字母（不区分大小写）作为集合S和T，然后计算他们的Jaccard相似性。

首先，我们需要确定集合S和T的元素。注意，每个字母只应计入一次，即使在名字中出现多次。

$S = \{a, l, i, c, e\}$

$T = \{e, l, s, a\}$

接着，我们需要确定S和T的交集( $S \cap T$ )和并集( $S \cup T$ )。

$S \cap T = \{a, l, e\}$ , 所以  $|S \cap T| = 3$

$S \cup T = \{a, l, i, c, e, s\}$ , 所以  $|S \cup T| = 6$

最后，我们将这两个数量插入到Jaccard相似性公式中：

Jaccard similarity =  $|S \cap T| / |S \cup T| = 3 / 6 = 0.5$

因此，名字 "Alice" 和 "Elsa" 的Jaccard相似性是 0.5。

## coding question

Create a python Spark program that does the following:

loads a text file (for example,  
<http://www.gutenberg.org/cache/epub/16328/pg16328.txt>)

finds and prints the length of the longest word (space delimited) in the file that contains the first letter of your first name and does not contain non-letter characters (i.e. the `isalpha()` method on a string in python should return True)

Full credit for efficient solutions that use Spark functions

下面是一个使用PySpark的程序，该程序加载一个文本文件，并找出包含我名字首字母且不含非字母字符的最长单词。

```
from pyspark import SparkConf, SparkContext
from operator import add

def filter_words(word):
    first_letter = 'c' # Assuming the first letter of my name is 'c'
    return word.startswith(first_letter) and word.isalpha()

def longest_word(word1, word2):
    return word1 if len(word1) > len(word2) else word2

# Initialize Spark
conf = SparkConf().setMaster("local").setAppName("LongestWord")
sc = SparkContext(conf = conf)

# Load the text file
rdd = sc.textFile("http://www.gutenberg.org/cache/epub/16328/pg16328.txt")

# Split the text file into words
words = rdd.flatMap(lambda line: line.split(" "))

# Filter words that start with the first letter of my name and contains only alphabets
filtered_words = words.filter(filter_words)

# Find the longest word
longest_word = filtered_words.reduce(longest_word)

print("Longest word: ", longest_word)
print("Length: ", len(longest_word))

sc.stop() # Stop Spark
```

这个程序首先加载一个文本文件，并将其分割为一个单词的RDD。然后，我们使用filter方法来找出所有以特定字母开头且只包含字母的单词。最后，我们使用reduce方法找出最长的单词。

请注意，运行此程序需要在你的环境中安装Spark和PySpark库，而且需要有访问该URL的权限。



# 必考概念

## 1. What are the benefits of computing with parallelism?

**Speed:** Parallel computing can significantly reduce the time it takes to compute complex tasks. By dividing a task among multiple processors that can operate simultaneously, you can complete the task much more quickly than if you used a single processor. This is particularly valuable for tasks that involve large amounts of data or complex calculations.

**Efficiency:** Parallel computing can be more efficient than sequential computing because it utilizes more of the computer's resources at the same time. While a single processor might spend a lot of time idle, multiple processors can all be used to work on different parts of the task at the same time.

**Problem Size:** Parallel computing allows for the tackling of larger problems that would be impractical or impossible to handle with a single processor. This is especially important in fields like scientific computing, where simulations and data analysis can involve enormous amounts of data.

**Cost-Effectiveness:** With the increasing availability and affordability of multi-core and multi-processor systems, parallel computing can be a cost-effective way to achieve high performance.

**Reliability:** In a distributed system, the failure of one machine does not halt the entire process as the workload can be picked up by other machines. This redundancy can lead to higher reliability.

**Concurrency:** It can handle many tasks simultaneously. Real-time systems with a large number of inputs and outputs or systems with a large number of users, like a web server, can benefit from this.

**Throughput:** For tasks that are not time-critical, using parallel computing to run them simultaneously can result in higher overall throughput. This can be a significant advantage in environments where many tasks need to be performed but each individual task is not particularly time-sensitive.